

Lab 3 -- Maze

Responsible TA: Jiwen Zhang

2022 11.15

Maze World Under MDP

- States: (col, row); cols*rows states in total

- Start state: State where agent starts the game
- State pair: Jump immediately to another state in a pair
- Goal state: End the game when agent reaches this state

- Actions

Up / Down / Left / Right;


Action result is determined

- Rewards

$n = \text{\#cells in this maze}$

$$R(s, a) = \begin{cases} -\frac{5}{n}, & \text{default reward} \\ -\frac{10}{n}, & \text{out of board / wall hitting} \\ 10, & \text{at the goal} \end{cases}$$

- Discount: 0.99



Maze World

- Problem:
 - Solve the Maze Problem based on MDP
 - Address: 10.192.9.230
- Methods for Solving MDP:
 - Value Iteration
 - Policy Iteration

Value Iteration

- Bellman Equation:

$$V^*(s) = \max_{a \in A(s)} R(s, a) + \gamma \sum_{s'} P(s'|s, a) * V^*(s')$$

- Synchronous Update

Use the state value of the last iteration as $V^*(s')$

- Asynchronous Update

Use the latest state value as $V^*(s')$

Policy Iteration

- Policy Evaluation

$$V^{\pi}(s) = R(s, \pi(s)) + \gamma \sum_{s'} P(s'|s, \pi(s)) * V^{\pi}(s')$$

- Policy Improvement

$$\pi(s) = \arg \max_{a \in A(s)} Q(s, a)$$

Code Template

- game.py: For play and visualization
 - Study mode : run your AI agent
 - Human mode : play by yourself
- maze_template.py:
 - Maze Environment
 - Your MazeRLAgent

Requirement for Lab

- Complete 3 functions in maze_template.py

jump directly to line 256 and start reading

- Line 396: policy_evaluation
- Line 412: policy_iteration
- Line 428: value_iteration

- Output: Print the **iteration numbers** and **optimal values** of all states using value iteration and policy iteration

Reminder

- Use **synchronous update** when you update state values.
- Use the given method in class **'MazeEnv'** to get legal actions of states and etc.

Pseudocode

- Value Iteration

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s'|s, a)$, rewards $R(s)$, discount γ .

ϵ , the accumulative error allowed in the utility of all states

local variables: U, U' , dict of utilities for states in S , initially zero

δ , the accumulative change in the utility of any stage in an iteration

repeat

$U \leftarrow U'; \delta \leftarrow 0$

for each state s **in** S **do**

$U'[s] \leftarrow \max_{a \in A(s)} R(s, a) + \gamma \sum_{s'} P(s'|s, a) U[s']$

$\delta \leftarrow \delta + |U'[s] - U[s]|$

until $\delta < \epsilon$

return U

Pseudocode

- Policy Iteration

function POLICY-ITERATION(*mdp*) **returns** a policy

inputs: *mdp*, an MDP with states S , actions $A(s)$, transition model $P(s'|s, a)$

local variables: U , a dict of utilities for states in S , initially zero

π , a dict of policy whose key is state, initially random

repeat

$U \leftarrow \text{POLICY_EVALUATION}(\pi, U, mdp)$

unchanged? \leftarrow true

for each state s **in** S **do**

if $\max_{a \in A(s)} Q(s, a) > Q(s, \pi[s])$ **then do**

$\pi[s] \leftarrow \arg \max_{a \in A(s)} Q(s, a)$

unchanged? \leftarrow false

until *unchanged?*

return π

Think by yourself :
How to compute the
Q-value in this problem?